

UE Apprentissage SIA S9

Supervise Learning

Mathieu FAUVEL

CESBIO, Université de Toulouse, CNES/CNRS/INRAe/IRD/UPS, Toulouse, FRANCE

October 11, 2023

Outline

Introduction

Introductory Example

- Linear model for regression

- Linear model for classification

Multi layer Perceptron

- Going deeper

- Needs some regularization

Convolutional Neural Networks

- Concepts

- How to improve the learning process

Recurrent Neural Networks

- When order matters

- Modern RNN

Outline

Introduction

Introductory Example

- Linear model for regression

- Linear model for classification

Multi layer Perceptron

- Going deeper

- Needs some regularization

Convolutional Neural Networks

- Concepts

- How to improve the learning process

Recurrent Neural Networks

- When order matters

- Modern RNN

- Contact: mathieu.fauvel@inrae.fr
- CV:

2004-2007 Ph.D. degree in Signal and Image Processing from the INP, Grenoble & the University of Iceland

2007-2008 Assistant Professor Grenoble

2008-2010 Post-doc position at INRIA - MISTIS Team

2010-2011 Assistant Professor Toulouse

2011-2018 Associate Professor at DYNAFOR & INP, Toulouse

Since 2018 Research (CRCN) at CESBIO, INRAe

- Research interests are:

Machine learning for environmental/ecological monitoring


Course Objectives

- Learn basics of modern machine learning
- Understand how each step works
 - ★ Data préparation
 - ★ Model definition
 - ★ Optimization step
- Implement various Deep Learning models in PyTorch
- Application to Computer Vision

What is Supervised (Machine) Learning ?

- **Artificial Intelligence**

Perform human tasks using computers and algorithms




What is Supervised (Machine) Learning ?

- **Artificial Intelligence**

Perform human tasks using computers and algorithms

- Mitchell, *The discipline of machine learning:*

Machine Learning is defined as the capacity of a computer program to improve its performance measure with observations



What is Supervised (Machine) Learning ?

- **Artificial Intelligence**


Perform human tasks using computers and algorithms

- Mitchell, *The discipline of machine learning:*

Machine Learning is defined as the capacity of a computer program to improve its performance measure with observations

- **Deep Learning**

- ★ Lot of data data
- ★ Complex model
- ★ High performance computing



Main Notations

Data

- Observed data $\mathbf{x} \in \mathbb{R}^d$, called *input variables* or *predictors*.
- Data to be predicted $y \in \mathbb{R}^p$, called *output variables* or *responses*.

Main Notations

Data

- Observed data $\mathbf{x} \in \mathbb{R}^d$, called *input variables* or *predictors*.
- Data to be predicted $y \in \mathbb{R}^p$, called *output variables* or *responses*.

Learning Problem

- If y are quantitative data: *regression* and $y \in \mathbb{R}^p$.
- If y are categorical data: *classification* and $y \in \{C_1, \dots, C_C\}$ with C_i refers to the i^{th} category.

Main Notations

Data

- Observed data $\mathbf{x} \in \mathbb{R}^d$, called *input variables* or *predictors*.
- Data to be predicted $y \in \mathbb{R}^p$, called *output variables* or *responses*.

Learning Problem

- If y are quantitative data: *regression* and $y \in \mathbb{R}^p$.
- If y are categorical data: *classification* and $y \in \{C_1, \dots, C_C\}$ with C_i refers to the i^{th} category.

Prediction function

$$\begin{aligned} f_{\theta} : \mathbb{R}^d &\rightarrow \mathbb{R}^p \\ \mathbf{x} &\mapsto y \end{aligned}$$

Online References

- Aston Zhang et al. “Dive into Deep Learning”. In: *arXiv preprint arXiv:2106.11342* (2021)
- Simon J.D. Prince. *Understanding Deep Learning*. MIT Press, 2023. URL: <https://udlbook.github.io/udlbook/>
- Sebastian Raschka, Yuxi (Hayden) Liu, and Vahid Mirjalili. *Machine Learning with PyTorch and Scikit-Learn*. Birmingham, UK: Packt Publishing, 2022. ISBN: 978-1801819312
- Kevin P. Murphy. *Probabilistic Machine Learning: An introduction*. MIT Press, 2022. URL: <http://probml.github.io/book1>
- Kevin P. Murphy. *Probabilistic Machine Learning: Advanced Topics*. MIT Press, 2023. URL: <http://probml.github.io/book2>

Outline

Introduction

Introductory Example

- Linear model for regression

- Linear model for classification

Multi layer Perceptron

- Going deeper

- Needs some regularization

Convolutional Neural Networks

- Concepts

- How to improve the learning process

Recurrent Neural Networks

- When order matters

- Modern RNN

Outline

Introduction

Introductory Example

- Linear model for regression

- Linear model for classification

Multi layer Perceptron

- Going deeper

- Needs some regularization

Convolutional Neural Networks

- Concepts

- How to improve the learning process

Recurrent Neural Networks

- When order matters

- Modern RNN

Univariate linear model

- $f(x) = wx + b$ and $\theta = (w, b)$
- Loss function: $\ell(f(x_i), y_i) = (f(x_i) - y_i)^2$
- Objective function:

$$\mathcal{L} = \frac{1}{n} \sum_{i=1}^n \ell_i$$

- Gradients update: $\theta^{t+1} = \theta^t - \eta \nabla_{\theta} \mathcal{L}$

$$\frac{\partial \mathcal{L}}{\partial w} = \frac{2}{n} \sum_{i=1}^n x_i (wx_i + b - y_i)$$

$$\frac{\partial \mathcal{L}}{\partial b} = \frac{2}{n} \sum_{i=1}^n (wx_i + b - y_i)$$

Univariate linear model

- $f(x) = wx + b$ and $\theta = (w, b)$
- Loss function: $\ell(f(x_i), y_i) = (f(x_i) - y_i)^2$
- Objective function:

$$\mathcal{L} = \frac{1}{n} \sum_{i=1}^n \ell_i$$

- Gradients update: $\theta^{t+1} = \theta^t - \eta \nabla_{\theta} \mathcal{L}$

$$\frac{\partial \mathcal{L}}{\partial w} = \frac{2}{n} \sum_{i=1}^n x_i (wx_i + b - y_i)$$

$$\frac{\partial \mathcal{L}}{\partial b} = \frac{2}{n} \sum_{i=1}^n (wx_i + b - y_i)$$

Work

Implement the 1D regression in pytorch of the following function:

$$f(x) = 2x - 1$$

Do the notebook:


[regression_1d_toy.ipynb](#)

Multivariate linear model

- $f(\mathbf{x}) = \mathbf{w}^\top \mathbf{x} + b$ and $\boldsymbol{\theta} = (\mathbf{w}, b)$
- Same loss and objective function and so same gradient updates ...

Multivariate linear model

- $f(\mathbf{x}) = \mathbf{w}^\top \mathbf{x} + b$ and $\boldsymbol{\theta} = (\mathbf{w}, b)$
- Same loss and objective function and so same gradient updates ...
- We can use other loss function




with H stands for *Hubert* loss:

$$H(f(\mathbf{x}_i) - y_i) = \begin{cases} \frac{1}{2}(f(\mathbf{x}_i) - y_i)^2 & \text{for } |f(\mathbf{x}_i) - y_i| \leq \delta, \\ (|f(\mathbf{x}_i) - y_i| - 0.5), & \text{otherwise.} \end{cases}$$

Multivariate linear model

- $f(\mathbf{x}) = \mathbf{w}^\top \mathbf{x} + b$ and $\boldsymbol{\theta} = (\mathbf{w}, b)$
- Same loss and objective function and so same gradient updates ...
- We can use other loss function



with H stands for *Hubert* loss:

$$H(f(\mathbf{x}_i) - y_i) = \begin{cases} \frac{1}{2}(f(\mathbf{x}_i) - y_i)^2 & \text{for } |f(\mathbf{x}_i) - y_i| \leq \delta, \\ (|f(\mathbf{x}_i) - y_i| - 0.5), & \text{otherwise.} \end{cases}$$

Work

On the **California housing** data set

- Prepare (train/validation) data
- Define a multi. linear model
- Optimize your hyperparameters (batch size, learning rate)
- Switch L2 to other loss function

Scaling feature for multivariate linear model with L2 loss function

- Suppose we have two features of different scale (e.g. because of different unit)

$$\mathbf{x}_1 \sim \mathcal{N}(0, 1) \text{ and } \mathbf{x}_2 \sim \mathcal{N}(10, 10)$$

- What is the impact on the gradient update ?

Scaling feature for multivariate linear model with L2 loss function

- Suppose we have two features of different scale (e.g. because of different unit)

$$\mathbf{x}_1 \sim \mathcal{N}(0, 1) \text{ and } \mathbf{x}_2 \sim \mathcal{N}(10, 10)$$

- What is the impact on the gradient update ?
- Noting $\mathbf{X} = [\mathbf{x}_1, \dots, \mathbf{x}_n]^\top$ and $\mathbf{y} = [y_1, \dots, y_n]^\top$

$$\nabla_{\mathbf{w}} = -\mathbf{X}^\top \underbrace{(\mathbf{y} - \mathbf{X}\mathbf{w})}_{\mathbf{e}} \Rightarrow \nabla_{\mathbf{w}_p} = -\sum_{i=1}^n \mathbf{e}_i \mathbf{x}_{ip}$$

- Parameter update

$$\mathbf{w}_p^{(t+1)} = \mathbf{w}_p^{(t)} + \frac{\eta}{n} \sum_{i=1}^n \mathbf{e}_i^{(t)} \mathbf{x}_{ip}$$

Scaling feature for multivariate linear model with L2 loss function

- Suppose we have two features of different scale (e.g. because of different unit)

$$\mathbf{x}_1 \sim \mathcal{N}(0, 1) \text{ and } \mathbf{x}_2 \sim \mathcal{N}(10, 10)$$

- What is the impact on the gradient update ?
- Noting $\mathbf{X} = [\mathbf{x}_1, \dots, \mathbf{x}_n]^\top$ and $\mathbf{y} = [y_1, \dots, y_n]^\top$

$$\nabla_{\mathbf{w}} = -\mathbf{X}^\top \underbrace{(\mathbf{y} - \mathbf{X}\mathbf{w})}_{\mathbf{e}} \Rightarrow \nabla_{\mathbf{w}_p} = -\sum_{i=1}^n \mathbf{e}_i \mathbf{x}_{ip}$$

- Parameter update

$$\mathbf{w}_p^{(t+1)} = \mathbf{w}_p^{(t)} + \frac{\eta}{n} \sum_{i=1}^n \mathbf{e}_i^{(t)} \mathbf{x}_{ip}$$

- Standardization of the input feature (and for regression the output too): [scaling_feature.ipynb](#)

$$\tilde{\mathbf{x}} = (\mathbf{x} - \boldsymbol{\mu}) \oslash \boldsymbol{\sigma}$$

Outline

Introduction

Introductory Example

Linear model for regression

Linear model for classification

Multi layer Perceptron

Going deeper

Needs some regularization

Convolutional Neural Networks

Concepts

How to improve the learning process

Recurrent Neural Networks

When order matters

Modern RNN

Classification as a regression problem

- Predict categorical variables ("Dogs", "Cats", "Monkey" ...)
- Estimation of the posterior probability

$$p(y_i = c | \mathbf{x}_i) \forall c \in \{1, \dots, C\} \text{ with } p(y_i = c | \mathbf{x}_i) \geq 0 \text{ and } \sum_{c=1}^C p(y_i = c | \mathbf{x}_i) = 1$$

- MAP rule: $\hat{c}_i = \arg \max_c p(y_i = c | \mathbf{x}_i)$

Classification as a regression problem

- Predict categorical variables ("Dogs", "Cats", "Monkey" ...)
- Estimation of the posterior probability

$$p(y_i = c | \mathbf{x}_i) \forall c \in \{1, \dots, C\} \text{ with } p(y_i = c | \mathbf{x}_i) \geq 0 \text{ and } \sum_{c=1}^C p(y_i = c | \mathbf{x}_i) = 1$$


- MAP rule: $\hat{c}_i = \arg \max_c p(y_i = c | \mathbf{x}_i)$
- Softmax linear regression

$$1 = \sum_{c=1}^C p(y_i = c | \mathbf{x}_i) = \sum_{c=1}^C \exp \left\{ \mathbf{w}_c^\top \mathbf{x}_i + b_c + K_i \right\} = \exp \{ K_i \} \sum_{c=1}^C \exp \left\{ \mathbf{w}_c^\top \mathbf{x}_i + b_c \right\}$$

thus



$$p(y_i = c | \mathbf{x}_i) = \frac{\exp \left\{ \mathbf{w}_c^\top \mathbf{x}_i + b_c \right\}}{\sum_{c'=1}^C \exp \left\{ \mathbf{w}_{c'}^\top \mathbf{x}_i + b_{c'} \right\}} = \frac{\exp(\mathbf{z}_{ic})}{\sum_{c'=1}^C \exp(\mathbf{z}_{ic'})} = \text{softmax}(\mathbf{z}_i)_c$$

Softmax



- Translation invariant: $\text{softmax}(\mathbf{z} + K) = \text{softmax}(\mathbf{z})$
- Preserve ordering relation: $\arg \max_c(\mathbf{p}_i) = \arg \max_c(\mathbf{z}_i)$
- If $\mathbf{z} \in \mathbb{R}^d$ with $d \rightarrow \infty$ then $\text{softmax}(\mathbf{z}) \rightarrow \text{"One-hot-vector"}$

Softmax



- Translation invariant: $\text{softmax}(\mathbf{z} + K) = \text{softmax}(\mathbf{z})$
- Preserve ordering relation: $\arg \max_c(\mathbf{p}_i) = \arg \max_c(\mathbf{z}_i)$
- If $\mathbf{z} \in \mathbb{R}^d$ with $d \rightarrow \infty$ then $\text{softmax}(\mathbf{z}) \rightarrow \text{"One-hot-vector"}$

Cross entropy loss function

- One hot encoding: $y_i = c \rightarrow \mathbf{y}_i \in \mathbb{R}^C$, $y_{ic} = 1$ if $i = c$ else 0
- Likelihood:

$$l(\mathbf{y}_i, \hat{\mathbf{y}}_i) = \prod_{c'=1}^C p(y_i = c' | \mathbf{x}_i)^{y_{ic'}} = \frac{\exp(\mathbf{z}_{ic})}{\sum_{c'=1}^C \exp(\mathbf{z}_{ic'})}$$


Cross entropy loss function

- One hot encoding: $y_i = c \rightarrow \mathbf{y}_i \in \mathbb{R}^C$, $y_{ic} = 1$ if $i = c$ else 0
- Likelihood:

$$l(\mathbf{y}_i, \hat{\mathbf{y}}_i) = \prod_{c'=1}^C p(y_i = c' | \mathbf{x}_i)^{y_{ic'}} = \frac{\exp(\mathbf{z}_{ic})}{\sum_{c'=1}^C \exp(\mathbf{z}_{ic'})}$$

- Negative log likelihood:

$$\begin{aligned}\ell(\mathbf{y}_i, \hat{\mathbf{y}}_i) &= - \sum_{c'=1}^C y_{ic'} \log \{p(y_i = c' | \mathbf{x}_i)\} \\ &= -\log \{p(y_i = c | \mathbf{x}_i)\} \\ &= -(\mathbf{w}_c^\top \mathbf{x}_i + b_c) + \underbrace{\log \left\{ \sum_{c'=1}^C \exp(\mathbf{w}_{c'}^\top \mathbf{x}_i + b_{c'}) \right\}}_{\text{logsumexp}(\mathbf{z}_i)}\end{aligned}$$



Work

What happens if:

- All $\mathbf{z}_{ic}$ are small ?
- One $\mathbf{z}_{ic}$ is very large ?
- All $\mathbf{z}_{ic}$ are very large ?

Multivariate linear classifier

- The model

$$\mathbf{p}_i = \text{softmax}(\mathbf{W}\mathbf{x}_i + \mathbf{b})$$

- Cross entropy loss function
- Same comments than for regression
 - ★ Split train, val and test
 - ★ Standardization
 - ★ Batch training
 - ★ Tune learning rate

Work

Implement a linear classifier in pytorch for the classification of **Fashion MNIST** data set.

Outline

Introduction

Introductory Example

- Linear model for regression

- Linear model for classification

Multi layer Perceptron

- Going deeper

- Needs some regularization

Convolutional Neural Networks

- Concepts

- How to improve the learning process

Recurrent Neural Networks

- When order matters

- Modern RNN

Outline

Introduction

Introductory Example

- Linear model for regression

- Linear model for classification

Multi layer Perceptron

- Going deeper

- Needs some regularization

Convolutional Neural Networks

- Concepts

- How to improve the learning process


Recurrent Neural Networks

- When order matters

- Modern RNN


Activation function

- As discussed earlier we need non-linearity between two layers
- Done with *activation function* between two layers
 - ★ Rectified Linear (ReLU): $\max(x, 0)$
 - ★ Sigmoid: $\frac{1}{1 + \exp(-x)}$
 - ★ Tanh: $\tanh(x)$



Multi Layer Perceptron (MLP)

- $\mathbf{x}^{(1)} = \sigma(\mathbf{W}^{(0)}\mathbf{x}^{(0)} + \mathbf{b}^{(0)})$
- $\mathbf{x}^{(2)} = \sigma(\mathbf{W}^{(1)}\mathbf{x}^{(1)} + \mathbf{b}^{(1)})$
- ...
- $\mathbf{y} = \mathbf{W}^H\mathbf{x}^H + \mathbf{b}^H$



$$\begin{aligned} &= \sigma(w_{1,0}x_0^{(0)} + w_{1,1}x_1^{(0)} + \dots + w_{1,d}x_d^{(0)} + b_1^{(0)}) \\ &= \sigma\left(\sum_{i=1}^d w_{1,i}x_i^{(0)} + b_1^{(0)}\right) \end{aligned}$$

$$\begin{pmatrix} x_1^{(1)} \\ x_2^{(1)} \\ \vdots \\ x_m^{(1)} \end{pmatrix} = \sigma \left[\begin{pmatrix} w_{1,0} & w_{1,1} & \dots & w_{1,d} \\ w_{2,0} & w_{2,1} & \dots & w_{2,d} \\ \vdots & \vdots & \ddots & \vdots \\ w_{m,0} & w_{m,1} & \dots & w_{m,d} \end{pmatrix} \begin{pmatrix} x_1^{(0)} \\ x_2^{(0)} \\ \vdots \\ x_d^{(0)} \end{pmatrix} + \begin{pmatrix} b_1^{(0)} \\ b_2^{(0)} \\ \vdots \\ b_m^{(0)} \end{pmatrix} \right]$$
$$\mathbf{x}^{(1)} = \sigma(\mathbf{W}^{(0)}\mathbf{x}^{(0)} + \mathbf{b}^{(0)})$$


From https://tikz.net/neural_networks/

Multi Layer Perceptron (MLP)

- $\mathbf{x}^{(1)} = \sigma(\mathbf{W}^{(0)}\mathbf{x}^{(0)} + \mathbf{b}^{(0)})$
- $\mathbf{x}^{(2)} = \sigma(\mathbf{W}^{(1)}\mathbf{x}^{(1)} + \mathbf{b}^{(1)})$
- ...
- $\mathbf{y} = \mathbf{W}^H\mathbf{x}^H + \mathbf{b}^H$

Work

- Why the last layer does not have non-linearity?
- For an univariate regression problem, shows that a 2-layers MLP is a piece-wise linear function.
- Implement an MLP classifier for Fashion MNIST data set



$$\begin{aligned} &= \sigma(w_{1,0}x_0^{(0)} + w_{1,1}x_1^{(0)} + \dots + w_{1,d}x_d^{(0)} + b_1^{(0)}) \\ &= \sigma\left(\sum_{l=1}^d w_{1,l}x_l^{(0)} + b_1^{(0)}\right) \end{aligned}$$

$$\begin{pmatrix} x_1^{(1)} \\ x_2^{(1)} \\ \vdots \\ x_m^{(1)} \end{pmatrix} = \sigma \left[\begin{pmatrix} w_{1,0} & w_{1,1} & \dots & w_{1,d} \\ w_{2,0} & w_{2,1} & \dots & w_{2,d} \\ \vdots & \vdots & \ddots & \vdots \\ w_{m,0} & w_{m,1} & \dots & w_{m,d} \end{pmatrix} \begin{pmatrix} x_1^{(0)} \\ x_2^{(0)} \\ \vdots \\ x_d^{(0)} \end{pmatrix} + \begin{pmatrix} b_1^{(0)} \\ b_2^{(0)} \\ \vdots \\ b_m^{(0)} \end{pmatrix} \right]$$
$$\mathbf{x}^{(1)} = \sigma(\mathbf{W}^{(0)}\mathbf{x}^{(0)} + \mathbf{b}^{(0)})$$

From https://tikz.net/neural_networks/

Outline

Introduction

Introductory Example

- Linear model for regression

- Linear model for classification

Multi layer Perceptron

- Going deeper

- Needs some regularization

Convolutional Neural Networks

- Concepts

- How to improve the learning process


Recurrent Neural Networks

- When order matters

- Modern RNN


Generalization Error

- *Generalization error*: Difference between the train accuracy and val/test accuracy
- *Overfitting*: High generalization error



Generalization Error


- *Generalization error*: Difference between the train accuracy and val/test accuracy
- *Overfitting*: High generalization error



Theory: favor simpler model and/or smooth model

Generalization Error

- *Generalization error*: Difference between the train accuracy and val/test accuracy
- *Overfitting*: High generalization error




Theory: favor simpler model and/or smooth model

- Early stopping: monitor the validation loss

Generalization Error

- *Generalization error*: Difference between the train accuracy and val/test accuracy
- *Overfitting*: High generalization error




Theory: favor simpler model and/or smooth model

- Early stopping: monitor the validation loss
- Weight decay: Tikhonov regularization

$$\mathcal{L} = \frac{1}{n} \sum_{i=1}^n \ell_i + \frac{\lambda}{H} \sum_{h=1}^H \|\mathbf{w}_h\|_2^2$$

Generalization Error

- *Generalization error*: Difference between the train accuracy and val/test accuracy
- *Overfitting*: High generalization error



Theory: favor simpler model and/or smooth model

- Early stopping: monitor the validation loss
- Weight decay: Tikhonov regularization
- Noise injection: Dropout

$$\mathcal{L} = \frac{1}{n} \sum_{i=1}^n \ell_i + \frac{\lambda}{H} \sum_{h=1}^H \|\mathbf{w}_h\|_2^2$$

Dropout

- Smoothness: $f(\mathbf{x} + \epsilon) \approx f(\mathbf{x})$
- Idea: Noise injection during the learning process such as $\mathbb{E}[\tilde{\mathbf{x}}] = \mathbf{x}$
- With MLP:

$$f(\mathbf{x}) = f_H \circ f_{H-1} \circ \dots \circ f_0(\mathbf{x})$$

$$\text{and } \mathbf{x}^{(h+1)} = f_h(\mathbf{x}^{(h)}) = \sigma(\mathbf{W}^{(h)}\mathbf{x}^{(h)} + \mathbf{b}^{(h)})$$

- Dropout: remove (set to zero) some internal features with probability p

$$\mathbf{x}^{(h+1)} = f_h(\tilde{\mathbf{x}}^{(h)}) = \sigma(\mathbf{W}^{(h)}(\mathbf{m}^{(h)} \odot \mathbf{x}^{(h)}) + \mathbf{b}^{(h)})$$

with

$$\mathbf{m}_i^{(h)} = \begin{cases} 0 & \text{with probability } p \\ 1/(1-p) & \text{with probability } 1-p \end{cases}$$

Dropout

- Smoothness: $f(\mathbf{x} + \epsilon) \approx f(\mathbf{x})$
- Idea: Noise injection during the learning process such as $\mathbb{E}[\tilde{\mathbf{x}}] = \mathbf{x}$
- With MLP:

$$f(\mathbf{x}) = f_H \circ f_{H-1} \circ \dots \circ f_0(\mathbf{x})$$

$$\text{and } \mathbf{x}^{(h+1)} = f_h(\mathbf{x}^{(h)}) = \sigma(\mathbf{W}^{(h)}\mathbf{x}^{(h)} + \mathbf{b}^{(h)})$$

- Dropout: remove (set to zero) some internal features with probability p

$$\mathbf{x}^{(h+1)} = f_h(\tilde{\mathbf{x}}^{(h)}) = \sigma(\mathbf{W}^{(h)}(\mathbf{m}^{(h)} \odot \mathbf{x}^{(h)}) + \mathbf{b}^{(h)})$$

with

$$\mathbf{m}_i^{(h)} = \begin{cases} 0 & \text{with probability } p \\ 1/(1-p) & \text{with probability } 1-p \end{cases}$$

- Enable during training and disable during inference (can be used to estimate posterior distribution)

$$f(\tilde{\mathbf{x}}) = f_H \circ D_H \circ f_{H-1} \circ \dots \circ D_1 \circ f_0(\mathbf{x})$$

Work

- Implement early stopping
- Implement dropout
- (Optional) Implement Tikhonov regularization (weight decay in pytorch)

Outline

Introduction

Introductory Example

- Linear model for regression

- Linear model for classification

Multi layer Perceptron

- Going deeper

- Needs some regularization

Convolutional Neural Networks

- Concepts

- How to improve the learning process

Recurrent Neural Networks

- When order matters

- Modern RNN

Outline

Introduction

Introductory Example

- Linear model for regression

- Linear model for classification

Multi layer Perceptron

- Going deeper

- Needs some regularization

Convolutional Neural Networks

Concepts


- How to improve the learning process

Recurrent Neural Networks

- When order matters

- Modern RNN

An image is not a vector of pixels



2D Convolution

- Fully connected layer:

$$x_{pq}^{(1)} = \sigma \left(\sum_{i,j=1}^{H,W} \mathbf{w}_{pqij}^{(0)} x_{ij}^{(0)} + b_{pq}^{(0)} \right)$$

- Act locally: restrict to neighborhood of pixel p, q

$$x_{pq}^{(1)} = \sigma \left(\sum_{i,j=-\Delta}^{i,j=+\Delta} \mathbf{w}_{ij}^{(0)} x_{(p+i)(q+j)}^{(0)} + b^{(0)} \right)$$

$$\begin{pmatrix} w_{pq11} & w_{pq12} & \dots & w_{pq1j} & \dots & w_{pq1W} \\ w_{pq21} & w_{pq22} & \dots & w_{pq2j} & \dots & w_{pq2W} \\ \vdots & \vdots & & \vdots & & \vdots \\ w_{pqi1} & w_{pqi2} & \dots & w_{pqij} & \dots & w_{pqiW} \\ \vdots & \vdots & & \vdots & & \vdots \\ w_{pqH1} & w_{pqH2} & \dots & w_{pqHj} & \dots & w_{pqHW} \end{pmatrix}$$

2D Convolution

- Fully connected layer:

$$x_{pq}^{(1)} = \sigma \left(\sum_{i,j=1}^{H,W} \mathbf{w}_{pqij}^{(0)} x_{ij}^{(0)} + b_{pq}^{(0)} \right)$$

- Act locally: restrict to neighborhood of pixel p, q

$$x_{pq}^{(1)} = \sigma \left(\sum_{i,j=-\Delta}^{i,j=+\Delta} \mathbf{w}_{ij}^{(0)} x_{(p+i)(q+j)}^{(0)} + b^{(0)} \right)$$

$$\begin{pmatrix} 0 & \dots & 0 & \dots & 0 \\ 0 & \dots & 0 & \dots & 0 \\ \vdots & w_{(i)(j-1)} & w_{(i-1)(j)} & w_{(i-1)(j+1)} & \vdots \\ 0 & w_{(i)(j-1)} & w_{(i)(j)} & w_{(i)(j+1)} & 0 \\ \vdots & w_{(i+1)(j-1)} & w_{(i+1)(j)} & w_{(i+1)(j+1)} & \vdots \\ 0 & \dots & 0 & \dots & 0 \end{pmatrix}$$

2D Convolution

- Fully connected layer:

$$x_{pq}^{(1)} = \sigma \left(\sum_{i,j=1}^{H,W} \mathbf{w}_{pqij}^{(0)} x_{ij}^{(0)} + b_{pq}^{(0)} \right)$$

- Act locally: restrict to neighborhood of pixel p, q

$$x_{pq}^{(1)} = \sigma \left(\sum_{i,j=-\Delta}^{i,j=+\Delta} \mathbf{w}_{ij}^{(0)} x_{(p+i)(q+j)}^{(0)} + b^{(0)} \right)$$

$$\begin{pmatrix} 0 & \dots & 0 & \dots & 0 \\ 0 & \dots & 0 & \dots & 0 \\ \vdots & w_{(i)(j-1)} & w_{(i-1)(j)} & w_{(i-1)(j+1)} & \vdots \\ 0 & w_{(i)(j-1)} & w_{(i)(j)} & w_{(i)(j+1)} & 0 \\ \vdots & w_{(i+1)(j-1)} & w_{(i+1)(j)} & w_{(i+1)(j+1)} & \vdots \\ 0 & \dots & 0 & \dots & 0 \end{pmatrix}$$

2D Convolution

- Fully connected layer:

$$x_{pq}^{(1)} = \sigma \left(\sum_{i,j=1}^{H,W} \mathbf{w}_{pqij}^{(0)} x_{ij}^{(0)} + b_{pq}^{(0)} \right)$$

- Act locally: restrict to neighborhood of pixel p, q

$$x_{pq}^{(1)} = \sigma \left(\sum_{i,j=-\Delta}^{i,j=+\Delta} \mathbf{w}_{ij}^{(0)} x_{(p+i)(j+q)}^{(0)} + b^{(0)} \right)$$

$$\begin{pmatrix} 0 & \dots & 0 & \dots & 0 \\ 0 & \dots & 0 & \dots & 0 \\ \vdots & w_{(i)(j-1)} & w_{(i-1)(j)} & w_{(i-1)(j+1)} & \vdots \\ 0 & w_{(i)(j-1)} & w_{(i)(j)} & w_{(i)(j+1)} & 0 \\ \vdots & w_{(i+1)(j-1)} & w_{(i+1)(j)} & w_{(i+1)(j+1)} & \vdots \\ 0 & \dots & 0 & \dots & 0 \end{pmatrix}$$

Work

Show that 2D convolution can be expressed as a linear layer.

Basics on 2D Convolution 1/2

Convolution

0	1	2	3	4
5	6	7	8	9
10	11	12	13	14
15	16	17	18	19

$\chi^{(h)}$

*

1	1	1
1	1	1
1	1	1

K

=

54	63	72
99	108	117

$\chi^{(h+1)}$

(More https://github.com/vdumoulin/conv_arithmetic/blob/master/README.md)

Basics on 2D Convolution 1/2

Convolution + Padding (1,1)

0	0	0	0	0	0	0
0	0	1	2	3	4	0
0	5	6	7	8	9	0
0	10	11	12	13	14	0
0	15	16	17	18	19	0
0	0	0	0	0	0	0

*

1	1	1
1	1	1
1	1	1

=

12	21	27	33	24
33	54	63	72	51
63	99	108	117	81
52	81	87	93	64

Basics on 2D Convolution 1/2

Convolution + Padding (1,1) + Stride (3, 2)

0	0	0	0	0	0	0
0	0	1	2	3	4	0
0	5	6	7	8	9	0
0	10	11	12	13	14	0
0	15	16	17	18	19	0
0	0	0	0	0	0	0

*

1	1	1
1	1	1
1	1	1

=

12	27	24
52	87	64

Basics on 2D Convolution 1/2

Convolution for multivalued images

0	1	2	3	4
5	6	7	8	9
10	11	12	13	14
15	16	17	18	19

*

1	1	1
1	1	1
1	1	1

=

54	63	72
99	108	117

+

=

288	279	270
243	234	225

19	18	17	16	15
14	13	12	11	10
9	8	7	6	5
4	3	2	1	0

*

2	2	2
2	2	2
2	2	2

=

234	216	198
144	126	108

Basics of 2D Convolution 2/2

- Multivariate input - Multivariate output

$$\mathbf{X}^{(h)} \in \mathbb{R}^{C_{(h)} \times H^{(h)} \times W^{(h)}}$$

$$\mathbf{K}^{(h)} \in \mathbb{R}^{C_{(h+1)} \times C_{(h)} \times k_H^h \times k_W^h}$$

$$\mathbf{X}^{(h+1)} \in \mathbb{R}^{C_{(h+1)} \times H^{(h+1)} \times W^{(h+1)}}$$

- Special case of 1×1 convolution: $(C_{(h+1)}, C_{(h)}, k_H, k_W) = (C_o, C_i, 1, 1)$
 - ★ Combine pixel across channels - no spatial operation
 - ★ Matrix product / linear layer in the channel dimension

$$\mathbf{X}^{(h+1)} = \text{reshape} \left\{ \sigma \left(\mathbf{K}^{(h)} \text{reshape} \{ \mathbf{X}^{(h)}, C_{(h)}, H^{(h)} W^{(h)} \} + \mathbf{b}^{(h)} \right), C_{(h+1)}, H^{(h)}, W^{(h)} \right\}$$

$$\text{with } \text{reshape} \{ \mathbf{X}^{(h)}, C_{(h)}, H^{(h)} W^{(h)} \} \in \mathbb{R}^{C_{(h)} \times H^{(h)} W^{(h)}}$$

Pooling

- Reduce the resolution of the image $\equiv$ Use larger convolution kernel
- Downsampling with simple operator on the neighborhood:
 - ★ Average value
 - ★ Max value (preferred)
- After the non-linearity

0	1	2	3	4
5	6	7	8	9
10	11	12	13	14
15	16	17	18	19

2×2 MaxPooling
Stride=2

6	8
16	18

Work

- Express multi-channels & multi-outputs convolution as a matrix product.
- Show that the successive convolution with two kernels is equivalent to one convolution.
- What stride size will you use for a max pooling with a kernel of size 2×2 ?
- Implement a CNN for the classification of *Fashion MNIST* data set: use one or two convolutional layers for the MLP.

Outline

Introduction

Introductory Example

- Linear model for regression

- Linear model for classification

Multi layer Perceptron

- Going deeper

- Needs some regularization

Convolutional Neural Networks

- Concepts

- How to improve the learning process

Recurrent Neural Networks

- When order matters

- Modern RNN

Batch Normalization

- Remember scaling problem in slide 10
- Similar problem may happen to each layer **but** we cannot use the same trick: the sample mean/variance change after each batch update !

Batch Normalization

- Remember scaling problem in slide 10
- Similar problem may happen to each layer **but** we cannot use the same trick: the sample mean/variance change after each batch update !
- Apply normalization per batch: *Batch Normalization*

$$\tilde{\mathbf{x}}_B = \gamma \odot \left[(\mathbf{x}_B - \boldsymbol{\mu}_B) \oslash \boldsymbol{\sigma}_B \right] + \boldsymbol{\beta}$$

with $\boldsymbol{\mu}_B$ and $\boldsymbol{\sigma}_B$ computed on batch B , and γ & $\boldsymbol{\beta}$ are learnable parameters

Batch Normalization

- Remember scaling problem in slide 10
- Similar problem may happen to each layer **but** we cannot use the same trick: the sample mean/variance change after each batch update !
- Apply normalization per batch: *Batch Normalization*

$$\tilde{\mathbf{x}}_B = \gamma \odot \left[(\mathbf{x}_B - \boldsymbol{\mu}_B) \oslash \boldsymbol{\sigma}_B \right] + \boldsymbol{\beta}$$


with $\boldsymbol{\mu}_B$ and $\boldsymbol{\sigma}_B$ computed on batch B , and γ & $\boldsymbol{\beta}$ are learnable parameters

Work

- What happen if the batch size is one ?
- What happen if the batch size is small ?
- What should we do in *prediction* mode ?
- Add normalization layer (2D) to your model.

Data Augmentation

- Noise injection to generate / augment data set: $(\mathbf{x}, y) \rightarrow (\tilde{\mathbf{x}}, y)$
- For image classification it is *relatively* easy:
 - ★ Geometric transformation: Rotation, translation, symmetry, radiometry ...
 - ★ Radiometric transformation: Invert, jitter ...
 - ★ *Any transformation that can model the variability of the acquisition process*



Work

Implement data augmentation with Pytorch for the training data only.

- Deep Neural Network: $f_h(\mathbf{x}^{(h)}) := f_h(\mathbf{x}^{(h)}, \theta_h)$

$$f(\mathbf{x}) = f_H \circ f_{H-1} \circ \dots \circ f_0(\mathbf{x})$$

Residual Network

- Deep Neural Network: $f_h(\mathbf{x}^{(h)}) := f_h(\mathbf{x}^{(h)}, \theta_h)$

$$f(\mathbf{x}) = f_H \circ f_{H-1} \circ \dots \circ f_0(\mathbf{x})$$

- Let $f_h = g_2 \circ g_1 \circ g_0$

$$\mathbf{x} \xrightarrow{g_0} \mathbf{x}^{(1)} \xrightarrow{g_1} \mathbf{x}^{(2)} \xrightarrow{g_2} f_h(\mathbf{x})$$

$$\frac{\partial f_h(\mathbf{x})}{\partial \alpha_0} = \underbrace{\frac{g_2}{g_1} \circ \frac{g_1}{g_0} \circ \frac{g_0}{\alpha_0}}_{\rightarrow}(\mathbf{x}) \quad \frac{\partial g(\mathbf{x})}{\partial \alpha_1} = \frac{g_2}{g_1} \circ \frac{g_1}{\alpha_1}(\mathbf{x}^{(1)}) \quad \frac{\partial g(\mathbf{x})}{\partial \alpha_2} = \frac{g_2}{\alpha_2}(\mathbf{x}^{(2)})$$

Residual Network

- Deep Neural Network: $f_h(\mathbf{x}^{(h)}) := f_h(\mathbf{x}^{(h)}, \theta_h)$

$$f(\mathbf{x}) = f_H \circ f_{H-1} \circ \dots \circ f_0(\mathbf{x})$$

- Let $f_h = g_2 \circ g_1 \circ g_0$ with *residual connection*

$$\mathbf{x} \xrightarrow{g_0} \oplus \rightarrow \mathbf{x}^{(1)} \xrightarrow{g_1} \oplus \rightarrow \mathbf{x}^{(2)} \xrightarrow{g_2} \oplus \rightarrow f_h(\mathbf{x})$$


$$\frac{\partial f_h(\mathbf{x})}{\partial \alpha_0} = \frac{g_2}{g_1} \circ \frac{g_1}{g_0} \circ \frac{g_0}{\alpha_0}(\mathbf{x}) + \frac{g_1}{g_0} \circ \frac{g_0}{\alpha_0}(\mathbf{x}) + \frac{g_0}{\alpha_0}(\mathbf{x})$$

- ★ Improve gradient propagation
- ★ Prevent vanishing/shattered gradient

Residual Network

- Deep Neural Network: $f_h(\mathbf{x}^{(h)}) := f_h(\mathbf{x}^{(h)}, \theta_h)$

$$f(\mathbf{x}) = f_H \circ f_{H-1} \circ \dots \circ f_0(\mathbf{x})$$



From Hao Li et al. “Visualizing the Loss Landscape of Neural Nets”. In: *Proceedings of the 32nd International Conference on Neural Information Processing Systems*. NIPS’18. Montréal, Canada: Curran Associates Inc., 2018, pp. 6391–6401

Outline

Introduction

Introductory Example

- Linear model for regression

- Linear model for classification

Multi layer Perceptron

- Going deeper

- Needs some regularization

Convolutional Neural Networks

- Concepts

- How to improve the learning process

Recurrent Neural Networks

- When order matters

- Modern RNN

Outline

Introduction

Introductory Example

- Linear model for regression

- Linear model for classification

Multi layer Perceptron

- Going deeper

- Needs some regularization

Convolutional Neural Networks

- Concepts

- How to improve the learning process


Recurrent Neural Networks

- When order matters

- Modern RNN

Sequential data

- Time series



- Text data

This course is really cool

Auto-regressive model

- Autogressive models: $p(\mathbf{x}_t | \mathbf{x}_{t-1}, \dots, \mathbf{x}_1)$
- Sequence model:

$$p(\mathbf{x}_T, \mathbf{x}_{T-1}, \dots, \mathbf{x}_1) = \prod_{t=2}^T p(\mathbf{x}_t | \mathbf{x}_{t-1}, \dots, \mathbf{x}_1) p(\mathbf{x}_1)$$

- Markov model:

$$p(\mathbf{x}_T | \mathbf{x}_{T-1}, \dots, \mathbf{x}_1) = \prod_{t=2}^T p(\mathbf{x}_t | \mathbf{x}_{t-1}) p(\mathbf{x}_1)$$

Auto-regressive model

- Autogressive models: $p(\mathbf{x}_t | \mathbf{x}_{t-1}, \dots, \mathbf{x}_1)$
- Sequence model:


$$p(\mathbf{x}_T, \mathbf{x}_{T-1}, \dots, \mathbf{x}_1) = \prod_{t=2}^T p(\mathbf{x}_t | \mathbf{x}_{t-1}, \dots, \mathbf{x}_1) p(\mathbf{x}_1)$$

- Markov model:

$$p(\mathbf{x}_T | \mathbf{x}_{T-1}, \dots, \mathbf{x}_1) = \prod_{t=2}^T p(\mathbf{x}_t | \mathbf{x}_{t-1}) p(\mathbf{x}_1)$$

- Example: $p(\mathbf{x}_t | \mathbf{x}_{t-1})$ could model
 - ★ The probability to get letter after a observing a specific letter (or a set of)
 - ★ The expected value of a signal given past values of the same signal

Latent autoregressive model



- $\mathbf{h}_t = f(\mathbf{x}_t, \mathbf{h}_{t-1}) = \sigma_h(\mathbf{w}_{hx}\mathbf{x}_t + \mathbf{w}_{hh}\mathbf{h}_{t-1} + \mathbf{b}_h)$
- $\mathbf{o}_t = g(\mathbf{h}_t) = \sigma_o(\mathbf{w}_{oh}\mathbf{h}_t + \mathbf{b}_o)$
- Backpropagation through time - BPTT

$$\ell = \frac{1}{T} \sum_{t=1}^T \ell(\mathbf{y}_t, \mathbf{o}_t)$$

with $\mathbf{h}_0 = \mathbf{0}$

Issue with BPTT

$$\begin{aligned}\frac{\partial \ell}{\partial \mathbf{W}_{hh}} &= \frac{1}{T} \sum_{t=1}^T \frac{\partial \ell(\mathbf{y}_t, \mathbf{o}_t)}{\partial \mathbf{W}_{hh}} \\&= \frac{1}{T} \sum_{t=1}^T \frac{\partial \ell(\mathbf{y}_t, \mathbf{o}_t)}{\partial \mathbf{o}_t} \frac{\partial \mathbf{o}_t}{\partial \mathbf{W}_{hh}} \\&= \frac{1}{T} \sum_{t=1}^T \frac{\partial \ell(\mathbf{y}_t, \mathbf{o}_t)}{\partial \mathbf{o}_t} \frac{\partial g(\mathbf{h}_t)}{\partial \mathbf{h}_t} \left[\frac{\partial f(\mathbf{x}_t, \mathbf{h}_{t-1})}{\partial \mathbf{W}_{hh}} + \frac{f(\mathbf{x}_t, \mathbf{h}_{t-1})}{\partial \mathbf{h}_{t-1}} \frac{\partial \mathbf{h}_{t-1}}{\partial \mathbf{W}_{hh}} \right] \\&= \frac{1}{T} \sum_{t=1}^T \frac{\partial \ell(\mathbf{y}_t, \mathbf{o}_t)}{\partial \mathbf{o}_t} \frac{\partial g(\mathbf{h}_t)}{\partial \mathbf{h}_t} \left[\frac{\partial f(\mathbf{x}_t, \mathbf{h}_{t-1})}{\partial \mathbf{W}_{hh}} + \sum_{p=1}^{t-1} \left(\prod_{q=p+1}^t \frac{\partial f(\mathbf{x}_q, \mathbf{h}_{q-1})}{\partial \mathbf{h}_{q-1}} \right) \frac{\partial f(\mathbf{x}_p, \mathbf{h}_{p-1})}{\partial \mathbf{W}_{hh}} \right]\end{aligned}$$

Outline

Introduction

Introductory Example

- Linear model for regression

- Linear model for classification

Multi layer Perceptron

- Going deeper

- Needs some regularization

Convolutional Neural Networks

- Concepts


- How to improve the learning process

Recurrent Neural Networks

- When order matters

- Modern RNN

Long Short Term Memory (LSTM)



- Forget gate

$$F_t = \sigma(\mathbf{w}_{fh} \mathbf{h}_{(t-1)} + \mathbf{w}_{fx} \mathbf{x}_t + \mathbf{b}_f)$$

- Input gate

$$I_t = \sigma(\mathbf{w}_{ih} \mathbf{h}_{(t-1)} + \mathbf{w}_{ix} \mathbf{x}_t + \mathbf{b}_i)$$

- Input node

$$\tilde{C}_t = \tanh(\mathbf{w}_{\tilde{c}h} \mathbf{h}_{(t-1)} + \mathbf{w}_{\tilde{c}x} \mathbf{x}_t + \mathbf{b}_{\tilde{c}})$$

- Output gate

$$O_t = \sigma(\mathbf{w}_{oh} \mathbf{h}_{(t-1)} + \mathbf{w}_{ox} \mathbf{x}_t + \mathbf{b}_o)$$

Good overview of LSTM: <https://colah.github.io/posts/2015-08-Understanding-LSTMs/>

Sequence prediction with RNN 1/2

Learning (toy) problem

Learn a function f that, given a fix sized sequence of characters, predicts the most probable next one from an alphabet

$$f(\text{This cour}) = \text{s}$$

$$f(\text{isawesom}) = \text{e}$$

Tokenization

Tokenization is the action of cutting input data into parts that can be embedded into a vector space.

```
text = "This course is awesome"
```

```
text_set = sorted(set(text))
```

```
[' ', 'T', 'a', 'c', 'e', 'h', 'i', 'm', 'o', 'r', 's', 'u', 'w']
```

```
char2int = {ch: i for i, ch in enumerate(text_set)}
```

```
{' ': 0, 'T': 1, 'a': 2, 'c': 3, 'e': 4, 'h': 5, 'i': 6, 'm': 7, 'o': 8,
```

```
↪ 'r': 9, 's': 10, 'u': 11, 'w': 12}
```

Sequence prediction with RNN 2/2

- Classification problem on sequence, $x_t = \text{token}$

$$\mathbf{x}_{T+1} = f(\mathbf{x}_T, \dots, \mathbf{x}_0)$$

- By-product: Once we have $p(\mathbf{x}_{T+1}|\mathbf{x}_T, \dots, \mathbf{x}_0)$ we can sample random text

$$p(\mathbf{x}_{T+1}|\mathbf{x}_T, \dots, \mathbf{x}_0) \sim \text{Multinomial}$$

Work

- Implement a simple tokenizer in pytorch
- Implement a RNN that learn to predict the next character of a sequence
- Once trained, generate some sentences of varying size

This work is licensed under a [Creative Commons “Attribution-ShareAlike 4.0 International”](#) license.

